

28/1/20

Unit-4 4. Transaction Management

Introduction to Transaction:

Collection of operations that perform a single logical unit of work are called transactions.

- A transaction is a unit of program execution that access and possibly update various data items.
- Transaction is initiated by a user program written in a high level data manipulation language (or) programming language [SQL, C++ (or) Java] where it is delimited by stmts of the form begin-transaction and end-transaction.
- The transaction consists of all operations executed between the begin and end transactions.
- The Transaction Management consists of 4 properties: these properties are often called as ACID properties.

- 1) Atomicity: Either all operations of the transaction are reflected properly in the database (or) None
- 2) Consistency: Execution of the transactions in isolation preserves the consistency of database
- 3) Isolation: Even though multiple transaction may execute concurrently, the system guarantees that each transaction is unaware of other transaction executing concurrently in the system

successfully the changes it has made to the database persist even if there are system failures.

→ Transaction access the data using two operations:

i) $\text{read}(x)$:- which transfers the data item x from database to a local buffer belonging to the transaction that executed the read operation.

ii) $\text{write}(x)$: Which transfers the data item x from the local buffer of the transaction that executed the write back to the database.

Example:

let T_1 be the transaction that transfers 50 rupees from Account A to account B.

This transaction can be defined as

$T_1 = /$
 $\text{read}(A);$

$A := A - 500;$

$\text{write}(A);$

$\text{read}(B);$

$B := B + 50;$

$\text{write}(B);$

i) consistency:-

The consistency requirement here is that the sum of A & B be unchanged by the execution of the transaction.

ensuring consistency for an individual transaction is the responsibility of the application programmer who codes the transaction.

g) Atomicity:

→ If the atomicity property is present all actions for the transactions are reflected in the database(s) none ..

→ The basic idea behind ensuring atomicity is the database system keeps track of the old values of any data on which a transaction performs a write. & if the transaction does not complete its execution, the database system restores the old values to make it appear as though the transaction never executed.

→ Ensuring atomicity is the responsibility of the database system itself and it is handled by the component called Transaction Management component.

Durability:-

The durability property guarantees that once a transaction completes successfully all the updates i.e, carried out on the database persist even if there is a system failure after the transaction completes execution.

→ we can guarantee durability by ensuring that either

1) the updates carried out by the transaction have been written to the disk before the transaction complete.

2) Information about the updates carried out by the transaction and written to disk is sufficient to enable the database to reconstruct the updates when the database system is restarted after the failure.

→ Ensuring durability is the responsibility of a software component of the database system called the Recovery management component.

Isolation:

The Isolation property of a transaction ensures that the concurrent execution of transaction results in a system state i.e., equivalent to a state that could have been obtained had these transactions executed one at a time in some order.

→ Ensuring the isolation property is the responsibility of a component of the database system called concurrency control component.

Transaction States:

A transaction must be in one of the following states active. The initial state the transaction stays in this state while

It is e:
partially c
has been
failed: A
executio
Aborted:
rolled b
restore
Start o
Commit
After E
state

1) It cc
only if
a real
error +
logic of
transa
transc

2) It c
thus
error
view:
or b
beca
fo

It is executing.
partially committed: After the final stmt has been executed.

failed: After the discovery that normal execution can no longer proceed.

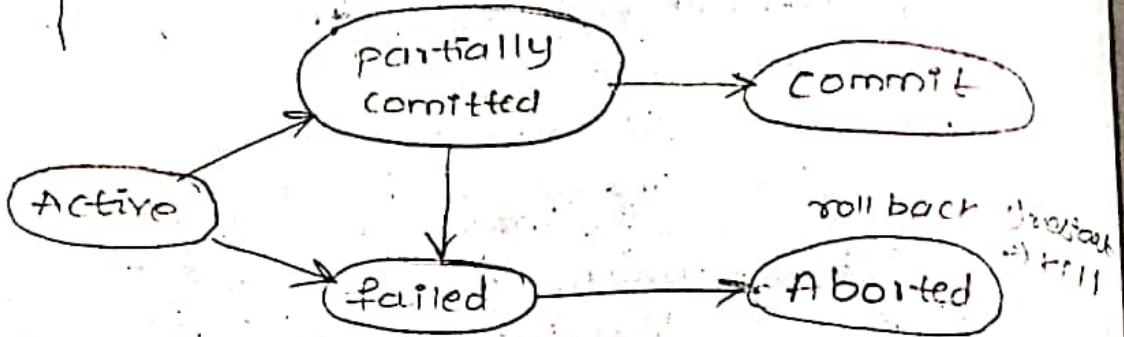
Aborted: After the transaction have been rolled back and database has been restored, to its state prior to the start of the transaction.

Committed :-

After successful completion at aborted state, the system has 2 options

- 1) It can restart the transaction but only if the transaction was aborted by a result of some hardware (or) software error that was not related to the internal logic of the transaction. A restarted transaction is considered to be a new transaction.
- 2) It can kill the transaction & usually thus so because of some internal logic error that can be corrected only by rewriting the application program or because the input was bad (a) because the desired data were not found in database.

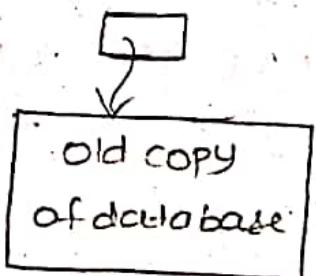
State diagram of a transaction



29/2/20

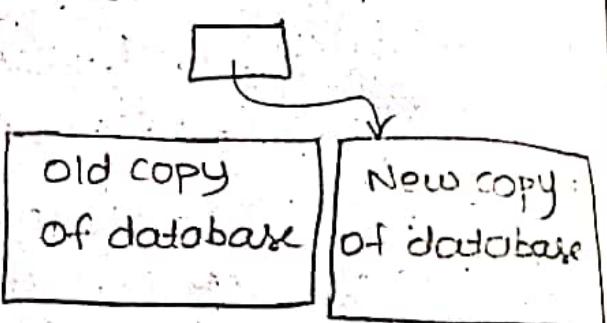
Implementation of Atomicity & Durability:

database pointer



a) Before update

database pointer



b) after update

fig :- shadow-copy technique of Atomicity & durability

→ The Recovery management component of db system can support atomicity & durability of a variety of schemes. A scheme called as the shadow-copy scheme. It is based on making copies of the database called shadow-copies.

→ A pointer called database pointer is maintained on disk it points to the current copy of the db.

→ The shadow wants to up copy of db dp copy.

→ If at any aborted the the old cop affected.

→ If the committed

- The OS is sure that the db has
- The data pointer t

The New copy of

then de have bec

the up

Transac

If tra point

the db

transac

COPY O

yster

→ If s update

→ The shadow-copy scheme is a transaction that wants to update the db. first creates a complete copy of db all updates are done on the new db copy.

→ If at any point the transaction has to be aborted the system deletes the new copy. The old copy of the db has not been affected.

→ If the transaction completes it is committed as follows.

1) The OS [operating system] is asked to make sure that all the pages of the new copy of the db have been written out to disk.

2) The database system updates the db pointer to point to the new copy of the db.

The New copy then becomes the current copy of the db. The old copy of the db is then deleted. The transaction is said to have been committed at the point where the updated db pointer is returned to disk.

Transaction failure:

If transaction fails at any time before db pointer is updated, the old content of the db are not effected, about the transaction by just deleting the new copy of db.

System failure:

→ If system fails at any time before the updated db pointer is returned to disk.

When the system restarts pt copy, read db pointer from the original contents of the db and none of the effects of the transaction will be visible on the db.

- 2) If the system fails after the db pointer has been updated on disk once new copy is updated on the disk its contents will not be damaged even if there is a system failure.

Thus the atomicity & durability properties of transactions are ensured by the shadow-copy implementation of the recovery management component.

Concurrent executions:-

The two good reasons for allowing concurrency

- 1) Improved throughput and resource utilization
- 2) Reduced waiting time

The database system must control the interaction among the concurrent transaction to prevent them from destroying the consistency of the db i.e., through a variety of mechanisms concurrency control scheme

Ex:- let T_1 & T_2 be two transactions that transfers funds from 1 account to another account.

transaction
account e
suppose th
are 1000 &

T_1 : read

A: A - 5

write C

read (1)

B: B + 5

write C

Serial

A Set

T_1 : rec

A: A

wait

read

B: E

wait

transaction T_1 transfers 50 rupees from account A to account B
suppose the current values of accounts A & B are 1000 & 2000 rupees respectively

T_1 : read(A);

A: A - 50;

write(A); // 950

read(B);

B: B + 50;

write(B); // 2050

T_2 : read(A) // 950

$$\text{temp} = A * 0.1; // 95 \\ A = A - \text{temp}; // 855$$

write(A); A = 855

read(B); // 2050

B: B + temp; // 2145

write(B); // 2145

Serial Schedule:-

A Serial schedule T_2 is followed by T_1 ,

T_1	T_2
T_1 : read(A); A: A - 50; write(A); read(B); B: B + 50; write(B);	T_2 : read(A) $\text{temp} = A * 0.1;$ A: A - temp; // 855 write(A); read(B); B: B + temp; // 2145 write(B); // 2145

Schedule-2 A serial schedule in which T_1 followed by T_2

T_1

T_1 : read(A); //1000
 $A := A - 50$; //1850
 write(A); //1850
 read(B); //2100
 $B := B + 50$; //2150
 write(B);

T_2

T_2 : read(A) //1000
 temp: $A * 0.1$; //100
 $A := A - \text{temp}$; //1900
 write(A); //1900
 read(B); //2050
 $B := B + \text{temp}$; //2150
 write(B); //2150

Schedule 1:

Read(A)

$A := A - 50$

write(A)

read(B)

$B := B + 50$

Schedule

T_1

Read(A)
 $A := A - 50$

Schedule-3: A concurrent schedule equivalent to Schedule-1

T_1

read(A) //1000
 $A := A - 50$; //950
 write(A); //950

 read(B)
 $B := B + 50$
 write(B)

T_2

Read(A) //950
 $\text{temp} := A * 0.1$ //95
 $A := A - \text{temp}$ //855
 write(A) //855

Read(B) //2050
 $B := B + \text{temp}$ //2145
 write(B) //2145

write(A)
 Read(B)
 $B := B + 50$
 write(B)

Schedule - 4: A concurrent schedule (Inconsistent)

T_1	T_2
Read(A) // 1000	
$A := A - 50 // 1050$	
	Read(A) // 1000
	$\text{temp} := A * 0.1 // 1100$
	$A := A - \text{temp} // 1050$
	Write(A) // 1050
Write(A) // 1050	
Read(B) // 2000	Read(B) // 2000
$B := B + 50 // 2050$	$B := B + \text{temp} // 2100$
	Write(B) // 2100

Schedule - 5: A concurrent schedule
(Inconsistent)

T_1	T_2
Read(A)	Read(A) // 1000
$A := A - 50$	$\text{temp} := A * 0.1 // 1100$
	$A := A - \text{temp} // 1050$
	Write(A) // 1050
Write(A) // 1050	Read(B) // 2000
Read(B) // 2000	$B := B + \text{temp} // 2100$
$B := B + 50 // 2050$	$B := B + 50 // 2050$
Write(B) // 2050	Write(B) // 2100

System carries out this task.

Concurrent schedule has to be always equivalent to a serial schedule.

** Serializability:

The database system must control concurrent execution of transactions to ensure that the database state remains consistent.

2 types of serializability.

- 1) conflict serializability
- 2) view serializability.
- 3) Conflict Serializability:

T_1	T_2
Read(A)	
Write(A)	
	Read(A)
	Write(A)
Read(B)	
Write(B)	
	Read(B)
	Write(B)

Let us consider a schedule A in which there are 2 consecutive instructions I_p and I_q of transactions T_1 & T_2 where $p \neq q$.

→ If I_i & I_j referred to different data item
then we can swap I_i & I_j without effecting
the results.
→ But I_i & I_j referred to the same data item
(Q) then the order of the 2 steps may matter
& cases to consider.

case-1: $I_i = \text{Read}(Q)$, $I_j = \text{Read}(Q)$

The order of I_i and I_j does not matter

case-2: $I_i = \text{Read}(Q)$, $I_j = \text{Write}(Q)$

If I_i comes before I_j then T_i does not
read the value of Q i.e., written by T_j
in instruction I_j . If I_j comes before I_i
then T_i reads the values of Q i.e. written
by T_j . thus the order of I_i & I_j matters

case-3: If $I_i = \text{Write}(Q)$ & $I_j = \text{Read}(Q)$

The order I_i & I_j matters for reasons
similar to that of previous case.

case-4: If $I_i = \text{Write}(Q)$, $I_j = \text{Write}(Q)$

The order of these instructions does not
effect either T_i or T_j but it will effect
the db which resides in disk so we
said that I_i and I_j conflict. If they
are operations by different transaction
on the data item.

9/3/20

Eoc:-

	T ₁	T ₂	T ₁	T ₂
	R(A)		R(A)	
	W(A)		W(A)	
		R(A) W(B)	R(B)	
	R(B)			W(A)
	W(B)			R(B) W(B)
		R(B) W(B)		

FST

	T ₁	T ₂	T ₁	T ₂
	R(A)		R(A)	
	W(A)		W(A)	
		R(A)	R(B)	
	R(B)			R(A)
	W(B)		W(B)	
		W(A)		W(A)
		R(B)		R(B)
		W(B)		RW(B)

V)

	T ₁	T ₂
	R(A)	
	W(A)	
	R(B)	
	W(B)	
		R(A)
		W(A)
		R(B)
		W(B)

the C
to i
schedu
es cont
thus
spnce

seia
conf
vpec

consi
same
bot

are
the
ij &

T₁
sch

gr
ve

2)
ex

pf-

sdg

t

E

the concept of conflict equivalence leads to the concept of conflict serializability. Schedule S is conflict serializable if it is conflict equivalent to a serial schedule. Thus schedule 0-3 is equivalent serializable since it is conflict equivalent to the serial schedule-1 that is called as conflict serializability.

View Serializability:

Consider two schedules S & S' where the same set of transactions participates in both schedules. The schedules S & S' are said to be view equivalent if three conditions are met.

- 1) For each data item $\langle Q \rangle$, if transaction T_j reads the initial value of $\langle Q \rangle$ in schedule S , then transaction T_i must, in schedule S' , also read the initial value of $\langle Q \rangle$.
- 2) For each data item $\langle Q \rangle$, if transaction T_j executes read $\langle Q \rangle$ in schedule S , and if that value was produced by a write $\langle Q \rangle$ operation executed by right $\langle Q \rangle$ operation executed by transaction T_i , then the read $\langle Q \rangle$ operation of transaction T_i must, in

Schedule s' , also read the value of Q that was produced by the write Ω operation of transaction T_j .

3) For each data item Q , the transaction that performs the final write Ω operation in schedule S , must perform the final write Ω operation in schedule S' .

→ In previous example schedule-1 is not view equivalent to schedule-2 but schedule-1 is view equivalent to schedule-3.

→ Schedule-S is said to be view serializable if it is view equivalent to Serial Schedule.

→ Every conflict serializable schedule is also view serializable^{but} there are view serializable schedules that are not conflict serializable.

Blind writes: In any transaction P if they perform write Ω operation without having read Ω operation, writes of this sort are called Blind writes. It occurs only in view-serializable schedule.

Implement for isolation to leave the db after it
→ no concurrent transaction
db before after it
→ no other accept it released
→ this transaction only see the goal is to f which can be viewed

Implementation of Isolation:

In Isolation property schedule must have to leave the database in consistent state & allow transaction failure to be handled in a save manner.

- To control the concurrency scheme a transaction acquires a lock on entire db before it starts and release the lock after it committed.
- No other transactions are allowed to access the db until the lock db has been released.
- This locking policy executes only one transaction at a time which means only serial schedules are generated.
- The goal of concurrency control scheme is to provide a high degree of consistency which ensuring that all schedules that can be generated are conflict or view serializable and cascadeless.

4/3/20

Recoverability:

Schedules are acceptable from the view point of consistency of the db if there are no transaction

If T_i fails we need to undo the effects of transaction T_i to ensure the atomicity property. If a system allows concurrent execution it is necessary to ensure any Transaction T_j depended on T_i where T_j Dead data written by T_i is also aborted.

1) Recoverable schedule:

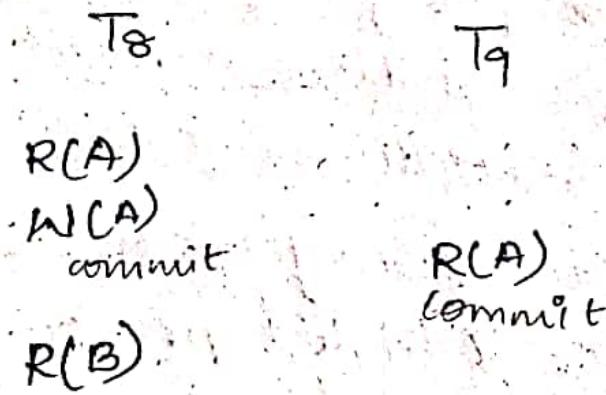


fig: schedule

→ If T_9 commit immediately after (executing) reading (A) instruction before T_8 commits, If T_8 fails then we have to abort T_9 to ensure atomicity but T_9 has already committed. where it

is impossible to recover correctly from the failure of T_8 . These type of schedules are known as Non-recoverable schedule.

- NO Non-recoverable is allow in the db
- A recoverable schedule is one where for each pair of transaction T_i & T_j such that T_j reads a data item previously written by T_i , the commit operation of T_i appears before the commit operation of T_j .

2) Cascadeless Schedule:

T_{10}	T_{11}	T_{12}
$R(A)$		
$R(B)$		
$w(A)$	$R(A)$	$R(A)$

Each and every transaction is dependent in this schedule. To recover correctly from transaction T_i we have to rollback several transactions. Then only the schedule is recoverable in which a single transaction failure leads to a series of transaction roll backs. This is called cascading roll back.

- If ignore cascading roll back occurs in the database cascadeless schedules are taken place.
- A cascadeless schedule is one where, for each pair of Transaction T_i & T_j , such that T_j reads a data item ~~be~~ written by T_i then the commit operation of T_i must appear before read operation of T_j .
- So every cascadeless schedule is also recoverable.

part-II

concurrency control:

When several transaction execute ^{concurrently} current in the db, the isolation property may no longer be preserved so interaction is necessary b/w the transaction this can be achieved by concurrency control Schemes:
lock based protocols:
to ensure serializability of transaction

9) Shared lock [s]: If a transaction T_j has obtained a shared mode lock S on data item Q, then T_j can read, but cannot write.

10) Exclusive lock [x]: If a transaction T_j has obtained a exclusive lock X on data item Q, then T_j can read & write



lock compatibility matrix component
[compatibility function]

T_1 : Transfer 50 from B to A.

T_2 : \rightarrow Read A & B values, display the sum of A & B values.

T_1	T_2
LOCK_X(B)	LOCK_S(A)
Read(B)	Read(A)
$B := B - 50$	unlock(A)
write(B)	LOCK_S(B)
UNLOCK(B)	R(B)
LOCK_X(A)	unlock(B)
$A := A + 50$	display(A+B)
W(A)	
UNLOCK(A)	

T ₁	T ₂
lock_x(B)	
r(B)	
B := B - 50	
w(B)	
unlock_(B)	lock_s(B)
B := B + 50	r(A)
lock_x(A)	unlock_(B)
r(A)	lock_s(A)
A := A + 50	r(A)
w(A)	unlock_s(A)
unlock_(A)	display(A+B)

→ Every transaction require a lock in an appropriate mode on data item(s).

the requirement is made to the concurrency control manager [CCM].

The transaction can proceed operation only after the CCM grants the lock to the transaction.

→ To access a data item transaction T_i must first lock that item. If the data item is already locked by another transaction in an incompatible mode, CCM will not grant the lock until all incompatible locks held by other transactions have been released.

so T_i has to wait.

T₁
Example:
lock
rec
B := B
w
unlock

lock
r
A :=
n
unlock

lock
B :=
unlock

lock
i
unlock

T₁
example:

T ₁	T ₂	CCM
lock-X(B) Read(B) B := B - 50 W(B) Unlock(B)	lock-SCA Read(A) Unlock(A) LOCK-SCA	→ Q-X(B, T ₁) → Q-SCA, T ₂) → Q-SCA, T ₂)
lock-X(CA) Read(CA) A := A + 50 write(CA) unlock(CA)	Read(B) unlock(B) Display(A+B)	Q-X(CA, T ₁)

T ₁	T ₂	CCM
LOCK-X(CB) Read(CB) B := B - 50 W(CB) Unlock(CB)	LOCK-SCB Read(CB) unlock(CB)	→ Q-X(CB, T ₁) → Q-SCB, T ₂) → Q-X(CA, T ₁)
LOCK-X(CA) Read(CA) A := A + 50 write(CA) unlock(CA)	lock-SCA Read(CA) unlock(CA) display(A+B)	→ Q-SCA, T ₁)

10.03.20

Deadlock:

Sometimes locking can be lead to an undesirable situation i.e., which leads to deadlock where T_1 is waiting for T_2 to release a lock on A and also T_2 is waiting for T_1 to release a lock on B where the transaction is stopped i.e., waiting for each other request is called as deadlock.

T_1	T_2
L-X(CB)	
R(CB)	
$B_i = B - 30$	
W(CB)	L-SC(A)
	R(CA)
	L-SC(B) X
X L-X(CA)	

fig: Deadlock Schedule.

To Overcome dead locks the system must rollback one of the two Transaction.

Roll back means release locks from transactions so that other transactions can execute there operations.

→ If we don't use locking (or) if we unlock data item's as soon as possible after reading (or) writing then the result is inconsistent data.

→ If we don't unlock a data item before request a lock on another data item deadlocks are going to occur.

→ To Handle the dead locks we are rolling back Transactions but the inconsistent states are not handled by the database system.

for all situations each transaction in system must follow a set of rules called locking protocol which indicates when to lock and unlock on data items.

② Granting of locks:

When a transaction request a lock on a data item in a particular mode and no other transaction has a lock on the same data item in a conflicting mode then the lock will be granted.

If T_2 has a shared lock on α , T_1 requests an exclusive mode lock on α then T_1 has to wait until T_2 releases the lock on α . If in the same time if T_3 requests a shared mode lock on α then the request is compatible so, the lock is granted to T_3 by T_2 instead of T_1 . If the same procedures continues the T_1 never gets the exclusive lock mode α so, T_1 is never executed. This process is known as starvation.

Avoid starvation of transaction by granting locks in the following manner:

When a transaction T_i request a lock on data-item α in a particular mode M the CCM grants the lock as

- i) There is no other transaction holding a lock α in a mode that conflicts with M .
- ii) There is no other transaction that is waiting for a lock on α and that made its lock request before T_i .
- iii) While performing the above manner a lock request never get blocked by a lock request that is made lock.

⑤ Two-phase locking protocol:

The two-phase locking protocol is one protocol that ensures serialisability.

Each transaction issues lock and unlock request in two phases.

1) Growing phase:

A Transaction may obtain locks, but may not release any lock is called a growing phase.

11) Shrinkage phase:

A Transaction may release locks but may not obtain any locks is called as shrinking phase.

iii) locking point (lock point):

The point in the schedule where the transaction has obtained its final lock i.e., end of growing phase.

- The transaction has to be ordered according to their locking points.
 - To be serializable the schedules have to be cascadeless but cascading rollback can occur under two-phase locking protocol.

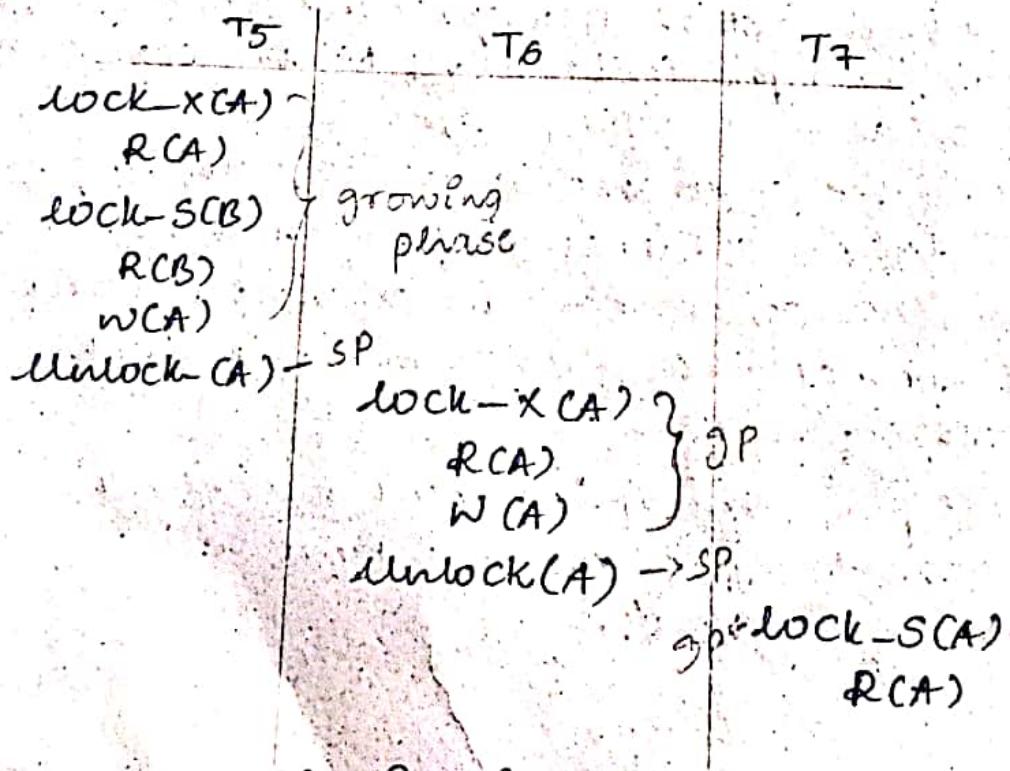


Fig. Partial Schedule Using
a phase locking.

It is when the cascade modification is called that

This requires but also by a trial

Commit
→ Ano
rígido de
seguir

→ Mo₂
Stric

\Rightarrow Refin

Wocka
Philpg
ij Don

is Upgr
ii) Down

14
growi
place

A S
approf
a do
misi

note
→ w/

Opera
follo

→ 60

If T_5 fails after T_3 reads data item (A) then the cascading rollback of T_6 & T_7 occurs. Cascading rollbacks can be avoided by a modification of two phase locking protocol called the strict two phase locking protocol.

This strict two-phase locking protocol requires not only that locking be two phase but also the all exclusive mode locks taken by a transaction be held until the txn. commits.

→ Another variant of 2-phase locking is the rigorous two phase locking protocol which requires that all locks be held until the transaction commits.

→ Most database system implements either strict (or) rigorous 2-phase locking.

→ Refinement of two phase locking protocol with the lock conversion:

lock conversions are of two types

- i) Upgrade
- ii) Downgrade.

i) Upgrade: A shared lock to an exclusive lock.

ii) Downgrade: An exclusive lock to a shared lock.

Upgrading can take place in only the growing phase, whereas downgrading can take place in only the shrinking phase.

A scheme automatically generates the appropriate lock and unlock instruction for a transaction, on the basis of read and write request from the transaction.

→ When a transaction T_i issues a read operation on Q, the system issues a lock-(Q)

followed by the read(Q) instruction.

→ When T_i issues a write(Q) operation, the system checks to see whether T_i already

holds a shared lock on α , if it does, then the system issues an upgrade(α) instruction followed by the write(α) instruction. Otherwise, the system issues a lock-x(α) instruction followed by the write(α) instruction.

→ All locks obtained by a transaction are unlocked after the transaction commits or aborts. (roll back).

④ Implementation of locking:-

→ A lock manager can be implemented as a process that receives messages from transaction and send messages in reply.

→ Lock Manager uses data structure i.e. linked list of records to maintain the request send by transaction.

→ Maintaining each and every request and order of a transaction with its name along with data item is a hashtable called lock table.
→ This algorithm guarantees that there is no starvation and ~~gives~~ uses shared memory instead of message passing for lock request (or) grant.

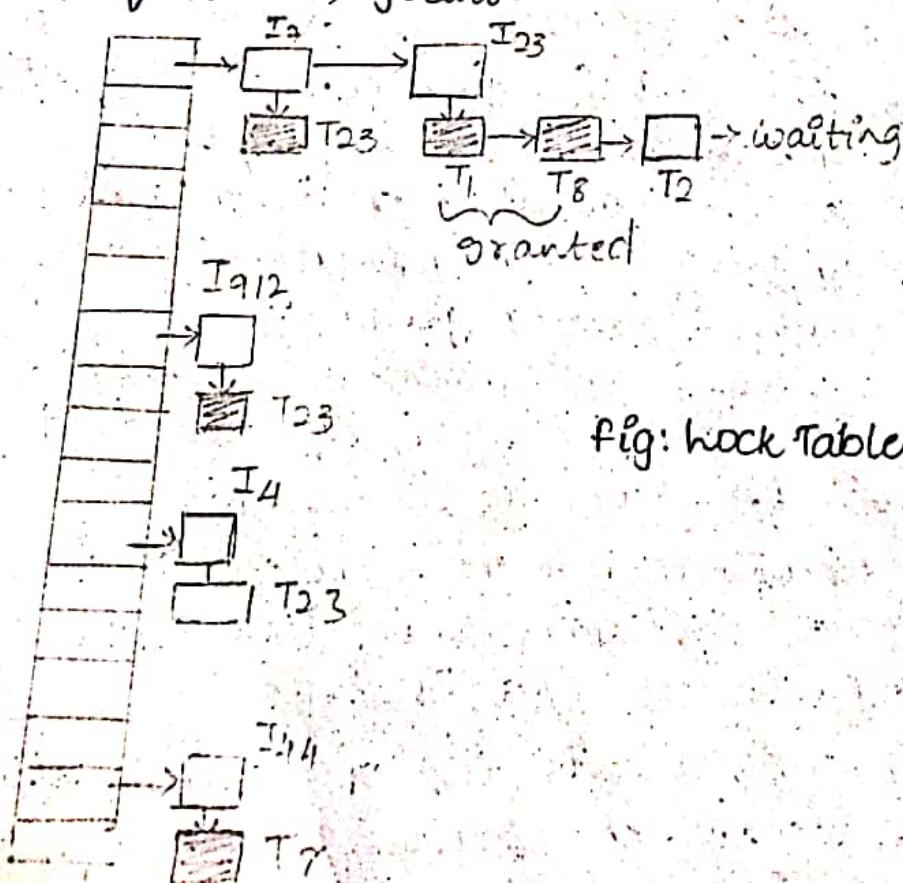


fig: lock Table

- the lock
1. When a record data item is created to the item if the grant request
2. When message record
3. If a transaction delete takes the

⑤ Gray

TO L
a pival
→ The P
may
graph
→ In
allo
can
close

i) the
ii) Such
by r
do

- the lock manager processes requests this way:
- 1. When a lock request message arrives, it adds a record to the end of the linked list for the data item. If the linked list is present otherwise it creates a new linked list. It always grants the first lock request on a data item but if the transaction request a lock on a data item on which a lock has already been granted, the lock manager grants the request only if it is compatible with earlier requests. Otherwise the request has to wait.
 - 2. When the lock manager receives an unlock messages from a transaction, it deletes the record for that data item in the linked list.
 - 3. If a transaction aborts, the lock manager deletes any waiting requests made by the transaction. Once the database system has taken appropriate actions to undo the transactions, it releases all locks held by the aborted transaction.

⑤ Graph Based Protocols:-

To construct locking protocols that are not 2 phase but ensure conflict serializability:

→ The partial ordering implies that the set 'D' may not be viewed as a directed acyclic graph called a database graph.

→ In the tree protocol, the only lock instruction allowed is lock exclusive each transaction T_i can lock a data item atmost once, and must observe the following rules.

- i) the first lock by T_i may be on any data item.
- ii) Subsequently a data item (a) can be locked by T_i only if the parent of a is currently locked by T_i .

- (iii) Data items may be unlocked at any time.
- (iv) A Data item that has been locked and unlocked by T_i cannot subsequently be relocked by T_i .

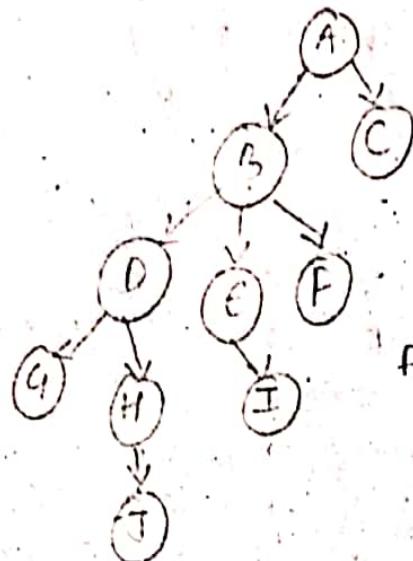


Fig: Tree Structured database graph.

→ This tree protocol does not ensure recoverability and cascadeless.

→ To ensure, the protocol can be modified not permit release of exclusive lock until the end the Transaction i.e; reduces concurrency.

** * Timestamp-Based Protocols:

In locking protocols the order b/w every pair of conflicting transaction is determined at execution time. By the first lock that they both request (i.e, incompatible modes).

An Ordering among transactions can be done in advance by using a common method that is timestamp-ordering scheme.

→ TimeStamp (TS):

The TS for each transaction T_i is assigned by the database system before the transaction T_i starts execution.

If T_j enters after TS assigned for T_i then $TS(T_i) < TS(T_j)$.

* Two Methods for Implementing this Scheme

i) Use the value of the system clock as the time stamp.

→ $TS(T_i)$ = value of the clock when the txn enters the system.

ii) Use a logical counter.

→ Incremented after a new time stamp has been assigned i.e., $TS(T_i)$ = value of the counter.

13/8/20

→ To implement this scheme, associated with each data item two timestamp values.

i) Write Time Stamp (w-Timestamp)(Q) : transaction denotes the largest time stamp of any time stamp that executed write(Θ) successfully.

ii) R-Timestamp(Q) : that denotes the largest time stamp of any transaction that executed read(Θ) successfully.

→ The timestamps are updated whenever a new Read(Θ) or write(Θ) instruction is executed.

* The timestamp ordering protocol:

It ensures that any conflicting read & write operations are executed in TS Order.

This protocol operates as follows:

1) Suppose that transaction T_i issues read(Θ)

a) If $TS_{OR}(T_i) \leq w-TS(\Theta)$ then T_i

needs to read a value of Θ that was already overwritten hence the read operation is rejected and T_i is roll back.

b) If $TS(T_i) > W-TS(Q)$ then the read operation is executed and $R-TS(Q)$ is set to the maximum of $R-TS(Q)$ & $TS(T_i)$

2) Suppose the transaction T_i issues write

a) If $TS(T_i) < R-TS(Q)$ then the value of Q that T_i is producing was needed previously and hence the value never be produced. Hence the system rejects the write operation and rolls T_i back.

b) If $TS(T_i) < R-TS(Q)$ then T_i is attempting to write an absolute value of Q . Hence the system rejects its write operation and rolls T_i back.

c) Otherwise, $[TS(T_i) \geq R-TS(Q) \text{ or } TS(T_i) \geq W-TS(Q)]$
then the system executes write operation and set $W-TS(Q)$ to $TS(T_i)$

Ex:- If there are two transactions T_{i4} & T_{i5} where T_{i4} is used to display the sum of acc A & acc B values while T_{i5} is used to transfer 50 Rs from B to A & display sum of A & B

display(A+B), R(B)

T15 : R(B)

B := B - 50

W(B)

R(A)

A := A + 50

W(A)

display(A+B)

R(A)

display(A+B)

R(B)

B := B - 50

W(B)

R(A)

A := A + 50

W(A)

display(A+B)

fig Schedule-3

Ex:-

T16	T17
R(Q)	
W(Q)	W(Q)

fig:- Schedule-4

* Thomas write Rule:

The modification to the time-stamp ordering protocol is called as Thomas write rule.

i) suppose the transaction T1 issues Read(Q).

a) If $TS(T_1) < TS(Q)$ then T1 needs to

read a value of Q that was already

written. Hence the read operation is

over written. Hence the read operation is rejected and T1 is roll back.

b) If $TS(T_i) > WL-TS(Q)$ then the read operation is executed and $R-TS(Q)$ is set to the minimum of $R-TS(Q)$ & $TS(T_i)$

2) Suppose the T_p issues $\text{write}(Q)$

a) If $TS(T_i) < R-TS(Q)$ then the value of Q that T_i is producing was needed previously the value never be produced. Hence the system rejects the write operation and rolls T_i back.

b) If $TS(T_i) < R-TS(Q)$ then T_i is attempting to write an absolute value of Q hence these write operation is rejected.

c) otherwise ($TS(T_i) \geq WL-TS(Q)$) the system executes the write operation and set $WL-TS(Q)$ to $TS(T_i)$.

NOTE:- $TS(T_i) \geq R-TS(Q)$ thomas write rule: ignore write operation

Validation based protocol:

To minimize the overhead of concurrency control scheme and to gain the knowledge of which transaction leads to conflict & whinot, A scheme is required to monitor the system.

there are 3 phases.

When the transaction T_i executes by system:

1) Read : The values of various data items are read, and are stored in variables local to T_i , All write operations are performed on temporary local variables without update on the actual database

2) Validation phase: It performs a validation test to determine whether the temporary local variables that holds the result of write operations can copy to database or not.

3) Write phase: If T_i succeeds in validation then the actual updates are applied to the database. Otherwise T_i is roll back

To perform the validation test, details of various phases of transaction T_i has to know.

1) $\text{Start}(T_i)$:- The time when T_i started its execution

2) $\text{Validation}(T_i)$:- The time when T_i finished pre-read phase and started its validation phase

3) $\text{finish}(T_i)$:- The time when T_i finished its write phase

Let $TS(T_i) = \text{Validation}(T_i)$

and $TS(T_i) < TS(T_j)$ where it means T_i appears before T_j

→ the validation test for transaction T_j requires that if transaction T_i with $TS(T_i) < TS(T_j)$ one of the following two conditions must hold

i) $\text{finish}(T_i) < \text{start}(T_j)$

Since T_i completes its execution before T_j started.

→ the serializability order is indeed maintained.

ii) the set of data items written by T_i does not intersect with the set of data items read by T_j , and at T_i complete its write phase before T_j starts its validation phase.

i.e. $\text{start}(T_j) < \text{finish}(T_i) &$

$\text{Validation}(T_j)$

This condition ensures that the $\text{write}(T_i)$ & T_j do not overlap.

R(A
Valid
displa

→ the
qua
pos

→ 10

P
Ce
exe
exe

R(B)

R(B)

B := B - 50

R(A)

A := A + 50

R(A)

<validate>

display(A+B)

<validate>

write(B)

write(A)

fig:- Schedule
produced by
using
validation

- the validation scheme automatically guards against cascading roll back but possibility of starvation.
- To avoid starvation, that validation phase is called Optimistic Concurrency Control Scheme, since transactions execute optimistically assume to finish execution and validate at end.