



# **Chapter 22: Object-Based Databases**

**Database System Concepts, 6th  
Ed.**

©Silberschatz, Korth and Sudarshan  
See [www.db-book.com](http://www.db-book.com) for conditions on re-use



# Chapter 22: Object-Based Databases

- n Complex Data Types and Object Orientation
- n Structured Data Types and Inheritance in SQL
- n Table Inheritance
- n Array and Multiset Types in SQL
- n Object Identity and Reference Types in SQL
- n Implementing O-R Features
- n Persistent Programming Languages
- n Comparison of Object-Oriented and Object-Relational Databases



# Object-Relational Data Models

- n Extend the relational data model by including object orientation and constructs to deal with added data types.
- n Allow attributes of tuples to have complex types, including non-atomic values such as nested relations.
- n Preserve relational foundations, in particular the declarative access to data, while extending modeling power.
- n Upward compatibility with existing relational languages.



# Complex Data Types

- n Motivation:
  - | Permit non-atomic domains (atomic  $\neq$  indivisible)
  - | Example of non-atomic domain: set of integers, or set of tuples
  - | Allows more intuitive modeling for applications with complex data
- n Intuitive definition:
  - | allow relations whenever we allow atomic (scalar) values
    - relations within relations
  - | Retains mathematical foundation of relational model
  - | Violates first normal form.





# Example of a Nested Relation

- n Example: library information system
- n Each book has
  - | title,
  - | a list (array) of authors,
  - | Publisher, with subfields *name* and *branch*, and
  - | a set of keywords
- n Non-1NF relation *books*

<i>title</i>	<i>author_array</i>	<i>publisher</i>	<i>keyword_set</i>
		( <i>name</i> , <i>branch</i> )	
Compilers	[Smith, Jones]	(McGraw-Hill, NewYork)	{parsing, analysis}
Networks	[Jones, Frick]	(Oxford, London)	{Internet, Web}



# 4NF Decomposition of Nested Relation

n Suppose for simplicity that title uniquely identifies a book

I In real world ISBN is a unique identifier

n Decompose *books* into 4NF using the schemas:

I  $(title, author, position)$

I  $(title, keyword)$

I  $(title, pub\_name, pub\_branch)$

n 4NF design requires users to include joins in their queries.

<i>title</i>	<i>author</i>	<i>position</i>
Compilers	Smith	1
Compilers	Jones	2
Networks	Jones	1
Networks	Frick	2

*authors*

<i>title</i>	<i>keyword</i>
Compilers	parsing
Compilers	analysis
Networks	Internet
Networks	Web

*keywords*

<i>title</i>	<i>pub_name</i>	<i>pub_branch</i>
Compilers	McGraw-Hill	New York
Networks	Oxford	London

*books4*



# Complex Types and SQL

- n Extensions introduced in SQL:1999 to support complex types:
  - | Collection and large object types
  - 4 Nested relations are an example of collection types
  - | Structured types
  - 4 Nested record structures like composite attributes
  - | Inheritance
  - | Object orientation
  - 4 Including object identifiers and references
- n Not fully implemented in any database system currently
  - | But some features are present in each of the major commercial database systems
  - 4 Read the manual of your database system to see what it supports



# Structured Types and Inheritance in SQL

**Structured types** (a.k.a. **user-defined types**) can be declared and used in SQL

**create type** *Name* **as**

```
(firstname      varchar(20),  
  lastname     varchar(20))  
final
```

**create type** *Address* **as**

```
(street        varchar(20),  
  city         varchar(20),  
  zipcode     varchar(20))
```

**not final**

Note: **final** and **not final** indicate whether subtypes can be created

Structured types can be used to create tables with composite attributes

```
create table person (  
  name      Name,  
  address   Address,  
  dateOfBirth date)
```

Dot notation used to reference components: *name.firstname*





# Structured Types (cont.)

n **User-defined row types**  
**create type** *PersonType* **as** (  
    *name* *Name*,  
    *address* *Address*,  
    *dateOfBirth* **date**)  
**not final**

n Can then create a table whose rows are a user-defined type  
**create table** *customer* **of** *CustomerType*

n Alternative using **unnamed row types**.

```
create table person_r(  
    name   row(firstname varchar(20),  
                    lastname varchar(20)),  
    address row(street   varchar(20),  
                    city     varchar(20),  
                    zipcode varchar(20)),  
    dateOfBirth date)
```



# Methods

- n Can add a method declaration with a structured type.

**method** *ageOnDate* (*onDate* **date**)

**returns** **interval year**

- n Method body is given separately.

**create instance method** *ageOnDate* (*onDate* **date**)

**returns** **interval year**

**for** *CustomerType*

**begin**

**return** *onDate* - **self.dateOfBirth**;

**end**

- n We can now find the age of each customer:

**select** *name.lastname*, *ageOnDate* (**current\_date**)

**from** *customer*



# Constructor Functions

n **Constructor functions** are used to create values of structured types

n E.g.

**create function** *Name*(*firstname* **varchar**(20), *lastname* **varchar**(20))

**returns** *Name*

**begin**

**set self.***firstname* = *firstname*;

**set self.***lastname* = *lastname*;

**end**

n To create a value of type *Name*, we use

**new** *Name*('John', 'Smith')

n Normally used in insert statements

**insert into** *Person* **values**

    (**new** *Name*('John', 'Smith'),

**new** *Address*('20 Main St', 'New York', '11001'),

**date** '1960-8-22');



n

# Type Inheritance

Suppose that we have the following type definition for people:

```
create type Person  
  (name varchar(20),  
   address varchar(20))
```

n

Using inheritance to define the student and teacher types

```
create type Student  
  under Person  
  (degree varchar(20),  
   department varchar(20))  
  
create type Teacher  
  under Person  
  (salary integer,  
   department varchar(20))
```

n

Subtypes can redefine methods by using **overriding method** in place of **method** in the method declaration



# Multiple Type Inheritance

- n SQL:1999 and SQL:2003 do not support multiple inheritance
- n If our type system supports multiple inheritance, we can define a type for teaching assistant as follows:
  - create type** *Teaching Assistant*  
**under** *Student, Teacher*
- n To avoid a conflict between the two occurrences of *department* we can rename them
  - create type** *Teaching Assistant*  
**under**  
*Student* **with** (*department as student\_dept* ),  
*Teacher* **with** (*department as teacher\_dept* )
- n Each value must have a **most-specific type**



# Table Inheritance

- n Tables created from subtypes can further be specified as **subtables**
- n E.g. **create table** *people* **of** *Person*;  
      **create table** *students* **of** *Student* **under** *people*;  
      **create table** *teachers* **of** *Teacher* **under** *people*;
- n Tuples added to a subtable are automatically visible to queries on the supertable
  - | E.g. query on *people* also sees *students* and *teachers*.
  - | Similarly updates/deletes on *people* also result in updates/deletes on subtables
  - | To override this behaviour, use “**only** *people*” in query
- n Conceptually, multiple inheritance is possible with tables
  - | e.g. *teaching\_assistants* under *students* and *teachers*
  - | *But is not supported in SQL currently*
- 4 So we cannot create a person (tuple in *people*) who is both a student and a teacher



# Consistency Requirements for Subtables

- n Consistency requirements on subtables and supertables.
- I Each tuple of the supertable (e.g. *people*) can correspond to at most one tuple in each of the subtables (e.g. *students* and *teachers*)
- I Additional constraint in SQL:1999:  
All tuples corresponding to each other (that is, with the same values for inherited attributes) must be derived from one tuple (inserted into one table).
- 4 That is, each entity must have a most specific type
- 4 We cannot have a tuple in *people* corresponding to a tuple each in *students* and *teachers*



# Array and Multiset Types in SQL

n Example of array and multiset declaration:

```
create type Publisher as
  (name          varchar(20),
   branch       varchar(20));
create type Book as
  (title         varchar(20),
   author_array varchar(20) array [10],
   pub_date      date,
   publisher     Publisher,
   keyword-set   varchar(20) multiset);
create table books of Book;
```





# Creation of Collection Values

- n Array construction  
**array** ['Silberschatz', `Korth`, `Sudarshan']
- n Multisets  
**multiset** ['computer', 'database', 'SQL']
- n To create a tuple of the type defined by the books relation:  
(`Compilers', **array**[`Smith', `Jones'],  
**new** *Publisher* (`McGraw-Hill', `New York'),  
**multiset** [`parsing', `analysis' ])
- n To insert the preceding tuple into the relation books  
**insert into** *books*  
**values**  
(`Compilers', **array**[`Smith', `Jones'],  
**new** *Publisher* (`McGraw-Hill', `New York'),  
**multiset** [`parsing', `analysis' ]);



# Querying Collection-Valued Attributes

n To find all books that have the word “database” as a keyword,

```
select title  
from books  
where ‘database’ in (unnest(keyword-set))
```

n We can access individual elements of an array by using indices

l E.g.: If we know that a particular book has three authors, we could write:

```
select author_array[1], author_array[2], author_array[3]  
from books  
where title = ‘Database System Concepts’
```

n To get a relation containing pairs of the form “title, author\_name” for each book and each author of the book

```
select B.title, A.author  
from books as B, unnest (B.author_array) as A (author)
```

n To retain ordering information we add a **with ordinality** clause

```
select B.title, A.author, A.position  
from books as B, unnest (B.author_array) with ordinality as  
A (author, position)
```



# Unnesting

n The transformation of a nested relation into a form with fewer (or no) relation-valued attributes is called **unnesting**.

n E.g.

```
select title, A as author, publisher.name as pub_name,  
        publisher.branch as pub_branch, K.keyword  
from books as B, unnest(B.author_array) as A (author),  
        unnest (B.keyword_set) as K (keyword)
```

n Result relation *flat\_books*

<i>title</i>	<i>author</i>	<i>pub_name</i>	<i>pub_branch</i>	<i>keyword</i>
Compilers	Smith	McGraw-Hill	New York	parsing
Compilers	Jones	McGraw-Hill	New York	parsing
Compilers	Smith	McGraw-Hill	New York	analysis
Compilers	Jones	McGraw-Hill	New York	analysis
Networks	Jones	Oxford	London	Internet
Networks	Frick	Oxford	London	Internet
Networks	Jones	Oxford	London	Web
Networks	Frick	Oxford	London	Web



# Nesting

- n **Nesting** is the opposite of unnesting, creating a collection-valued attribute
- n Nesting can be done in a manner similar to aggregation, but using the function **collect()** in place of an aggregation operation, to create a multiset
- n To nest the *flat\_books* relation on the attribute *keyword*:

```
select title, author, Publisher (pub_name, pub_branch ) as publisher,  
        collect (keyword) as keyword_set  
from flat_books  
groupby title, author, publisher
```

- n To nest on both authors and keywords:

```
select title, collect (author ) as author_set,  
        Publisher (pub_name, pub_branch) as publisher,  
        collect (keyword ) as keyword_set  
from flat_books  
group by title, publisher
```



## Nesting (Cont.)

- n Another approach to creating nested relations is to use subqueries in the **select** clause, starting from the 4NF relation *books4*

```
select title,  
      array (select author  
             from authors as A  
             where A.title = B.title  
             A.position) as author_array,  
      Publisher (pub-name, pub-branch) as publisher,  
      multiset (select keyword  
               from keywords as K  
               where K.title = B.title) as keyword_set  
from books4 as B
```

**order by**



# Object-Identity and Reference Types

- n Define a type *Department* with a field *name* and a field *head* which is a reference to the type *Person*, with table *people* as scope:  
**create type** *Department* (  
    *name* **varchar** (20),  
    *head* **ref** (*Person*) **scope** *people*)
- n We can then create a table *departments* as follows  
**create table** *departments* **of** *Department*
- n We can omit the declaration **scope** *people* from the type declaration and instead make an addition to the **create table** statement:  
**create table** *departments* **of** *Department*  
    (*head* **with options** **scope** *people*)
- n Referenced table must have an attribute that stores the identifier, called the **self-referential attribute**  
**create table** *people* **of** *Person*  
    **ref is** *person\_id* **system generated;**



# Initializing Reference-Typed Values

- n To create a tuple with a reference value, we can first create the tuple with a null reference and then set the reference separately:

```
insert into departments
values ('CS', null)
update departments
set head = (select p.person_id
            from people as p
            where name = 'John')
where name = 'CS'
```



# User Generated Identifiers

- n The type of the object-identifier must be specified as part of the type definition of the referenced table, and
- n The table definition must specify that the reference is user generated

```
create type Person  
  (name varchar(20)  
   address varchar(20))  
  ref using varchar(20)  
create table people of Person  
  ref is person_id user generated
```

- n When creating a tuple, we must provide a unique value for the identifier:

```
insert into people (person_id, name, address) values  
  ('01284567', 'John', '23 Coyote Run')
```

- n We can then use the identifier value when inserting a tuple into *departments*

- l Avoids need for a separate query to retrieve the identifier:

```
insert into departments  
  values('CS', '02184567')
```





# User Generated Identifiers (Cont.)

- n Can use an existing primary key value as the identifier:

```
create type Person  
  (name varchar (20) primary key,  
   address varchar(20))  
ref from (name)  
create table people of Person  
  ref is person_id derived
```

- n When inserting a tuple for *departments*, we can then use

```
insert into departments  
  values(`CS`,`John`)
```



# Path Expressions

- n Find the names and addresses of the heads of all departments:  
**select** *head*  $\rightarrow$  *name*, *head*  $\rightarrow$  *address*  
**from** *departments*
- n An expression such as “*head*  $\rightarrow$  *name*” is called a **path expression**
- n Path expressions help avoid explicit joins
  - | If department head were not a reference, a join of *departments* with *people* would be required to get at the address
  - | Makes expressing the query much easier for the user



# Implementing O-R Features

- n Similar to how E-R features are mapped onto relation schemas
- n Subtable implementation
  - | Each table stores primary key and those attributes defined in that table
  - or,
  - | Each table stores both locally defined and inherited attributes



# Persistent Programming Languages

- n Languages extended with constructs to handle persistent data
- n Programmer can manipulate persistent data directly
  - | no need to fetch it into memory and store it back to disk (unlike embedded SQL)
- n Persistent objects:
  - | **Persistence by class** - explicit declaration of persistence
  - | **Persistence by creation** - special syntax to create persistent objects
  - | **Persistence by marking** - make objects persistent after creation
  - | **Persistence by reachability** - object is persistent if it is declared explicitly to be so or is reachable from a persistent object



# Object Identity and Pointers

- n Degrees of permanence of object identity
  - | **Intraprocedure**: only during execution of a single procedure
  - | **Intraprogram**: only during execution of a single program or query
  - | **Interprogram**: across program executions, but not if data-storage format on disk changes
  - | **Persistent**: interprogram, plus persistent across data reorganizations
- n Persistent versions of C++ and Java have been implemented
  - | C++
    - 4 ODMG C++
    - 4 ObjectStore
  - | Java
    - 4 Java Database Objects (JDO)



# Persistent C++ Systems

- n Extensions of C++ language to support persistent storage of objects
- n Several proposals, ODMG standard proposed, but not much action of late
  - | **persistent pointers**: e.g. `d_Ref<T>`
  - | **creation of persistent objects**: e.g. `new (db) T()`
  - | **Class extents**: access to all persistent objects of a particular class
  - | **Relationships**: Represented by pointers stored in related objects
  - 4 Issue: consistency of pointers
  - 4 Solution: extension to type system to automatically maintain back-references
  - | **Iterator interface**
  - | **Transactions**
  - | **Updates**: `mark_modified()` function to tell system that a persistent object that was fetched into memory has been updated
  - | **Query language**



# Persistent Java Systems

- n Standard for adding persistence to Java : **Java Database Objects (JDO)**
  - | Persistence by reachability
  - | Byte code enhancement
    - 4 Classes separately declared as persistent
    - 4 Byte code modifier program modifies class byte code to support persistence
      - E.g. Fetch object on demand
      - Mark modified objects to be written back to database
  - | Database mapping
    - 4 Allows objects to be stored in a relational database
  - | Class extents
  - | Single reference type
    - 4 no difference between in-memory pointer and persistent pointer
    - 4 Implementation technique based on **hollow objects** (a.k.a. **pointer swizzling**)



# Object-Relational Mapping

- n **Object-Relational Mapping (ORM)** systems built on top of traditional relational databases
- n Implementor provides a mapping from objects to relations
  - | Objects are purely transient, no permanent object identity
- n Objects can be retried from database
  - | System uses mapping to fetch relevant data from relations and construct objects
  - | Updated objects are stored back in database by generating corresponding update/insert/delete statements
- n The **Hibernate** ORM system is widely used
  - | described in Section 9.4.2
  - | Provides API to start/end transactions, fetch objects, etc
  - | Provides query language operating directly on object model
  - 4 queries translated to SQL
- n Limitations: overheads, especially for bulk updates





# Comparison of O-O and O-R Databases

- n **Relational systems**
  - l simple data types, powerful query languages, high protection.
- n **Persistent-programming-language-based OODBs**
  - l complex data types, integration with programming language, high performance.
- n **Object-relational systems**
  - l complex data types, powerful query languages, high protection.
- n **Object-relational mapping systems**
  - l complex data types integrated with programming language, but built as a layer on top of a relational database system
- n Note: Many real systems blur these boundaries
  - l E.g. persistent programming language built as a wrapper on a relational database offers first two benefits, but may have poor performance.



# **End of Chapter 22**

**Database System Concepts, 6th  
Ed.**

©Silberschatz, Korth and Sudarshan  
See [www.db-book.com](http://www.db-book.com) for conditions on re-use



## Figure 22.05

<i>instructor</i>	
<u><i>ID</i></u>	
<i>name</i>	
<i>first_name</i>	
<i>middle_initial</i>	
<i>last_name</i>	
<i>address</i>	
<i>street</i>	
<i>street_number</i>	
<i>street_name</i>	
<i>apt_number</i>	
<i>city</i>	
<i>state</i>	
<i>zip</i>	
{ <i>phone_number</i> }	
<i>date_of_birth</i>	
<i>age</i> ( )	



**Figure 22.07**

